



Host setup

NetApp Solutions

Dorian Henderson, Ivana Devine
July 22, 2021

Table of Contents

- Host setup 1
 - Configure multipathing 1
 - Create LVM volume groups and logical volumes 6
 - Create file systems 7
 - Create mount points 8
 - Mount file systems 9

Host setup

Previous: [SAP HANA storage connector API](#).

Before setting up the host, NetApp SAN Host Utilities must be downloaded from the [NetApp Support](#) site and installed on the HANA servers. The Host Utility documentation includes information about additional software that must be installed depending on the FCP HBA used.

The documentation also contains information about multipath configurations that are specific to the Linux version used. This document covers the required configuration steps for SLES 15 and Red Hat Enterprise Linux 7.6 or higher, as described in the [Linux Host Utilities 7.1 Installation and Setup Guide](#).

Configure multipathing



Steps 1 to 6 must be performed on all worker and standby hosts in the SAP HANA multiple-host configuration.

To configure multipathing, complete the following steps:

1. Run the Linux `rescan-scsi-bus.sh -a` command on each server to discover new LUNs.
2. Run the `sanlun lun show` command and verify that all required LUNs are visible. The following example shows the `sanlun lun show` command output for a 2+1 multiple-host HANA system with two data LUNs and two log LUNs. The output shows the LUNs and the corresponding device files, such as LUN `SS3_data_mnt00001` and the device file `/dev/sdag`. Each LUN has eight FC paths from the host to the storage controllers.

```
stlrx300s8-6:~ # sanlun lun show
controller(7mode/E-Series)/
device          host      lun
vserver(cDOT/FlashRay)    lun-pathname
filename        adapter  protocol  size    product
-----
hana            /vol/SS3_log_mnt00002/SS3_log_mnt00002
/dev/sdah       host11    FCP       512.0g  cDOT
hana            /vol/SS3_data_mnt00001/SS3_data_mnt00001
/dev/sdag       host11    FCP       1.2t    cDOT
hana            /vol/SS3_data_mnt00002/SS3_data_mnt00002
/dev/sdaf       host11    FCP       1.2t    cDOT
hana            /vol/SS3_log_mnt00002/SS3_log_mnt00002
/dev/sdae       host11    FCP       512.0g  cDOT
hana            /vol/SS3_data_mnt00001/SS3_data_mnt00001
/dev/sdad       host11    FCP       1.2t    cDOT
hana            /vol/SS3_data_mnt00002/SS3_data_mnt00002
/dev/sdac       host11    FCP       1.2t    cDOT
hana            /vol/SS3_log_mnt00002/SS3_log_mnt00002
/dev/sdab       host11    FCP       512.0g  cDOT
```

hana			/vol/SS3_data_mnt00001/SS3_data_mnt00001
/dev/sdaa	host11	FCP	1.2t cDOT
hana			/vol/SS3_data_mnt00002/SS3_data_mnt00002
/dev/sdz	host11	FCP	1.2t cDOT
hana			/vol/SS3_log_mnt00002/SS3_log_mnt00002
/dev/sdy	host11	FCP	512.0g cDOT
hana			/vol/SS3_data_mnt00001/SS3_data_mnt00001
/dev/sdx	host11	FCP	1.2t cDOT
hana			/vol/SS3_data_mnt00002/SS3_data_mnt00002
/dev/sdw	host11	FCP	1.2t cDOT
hana			/vol/SS3_log_mnt00001/SS3_log_mnt00001
/dev/sdv	host11	FCP	512.0g cDOT
hana			/vol/SS3_log_mnt00001/SS3_log_mnt00001
/dev/sdu	host11	FCP	512.0g cDOT
hana			/vol/SS3_log_mnt00001/SS3_log_mnt00001
/dev/sdt	host11	FCP	512.0g cDOT
hana			/vol/SS3_log_mnt00001/SS3_log_mnt00001
/dev/sds	host11	FCP	512.0g cDOT
hana			/vol/SS3_log_mnt00002/SS3_log_mnt00002
/dev/sdr	host10	FCP	512.0g cDOT
hana			/vol/SS3_data_mnt00001/SS3_data_mnt00001
/dev/sdq	host10	FCP	1.2t cDOT
hana			/vol/SS3_data_mnt00002/SS3_data_mnt00002
/dev/sdp	host10	FCP	1.2t cDOT
hana			/vol/SS3_log_mnt00002/SS3_log_mnt00002
/dev/sdo	host10	FCP	512.0g cDOT
hana			/vol/SS3_data_mnt00001/SS3_data_mnt00001
/dev/sdn	host10	FCP	1.2t cDOT
hana			/vol/SS3_data_mnt00002/SS3_data_mnt00002
/dev/sdm	host10	FCP	1.2t cDOT
hana			/vol/SS3_log_mnt00002/SS3_log_mnt00002
/dev/sdl	host10	FCP	512.0g cDOT
hana			/vol/SS3_data_mnt00001/SS3_data_mnt00001
/dev/sdk	host10	FCP	1.2t cDOT
hana			/vol/SS3_data_mnt00002/SS3_data_mnt00002
/dev/sdj	host10	FCP	1.2t cDOT
hana			/vol/SS3_log_mnt00002/SS3_log_mnt00002
/dev/sdi	host10	FCP	512.0g cDOT
hana			/vol/SS3_data_mnt00001/SS3_data_mnt00001
/dev/sdh	host10	FCP	1.2t cDOT
hana			/vol/SS3_data_mnt00002/SS3_data_mnt00002
/dev/sdg	host10	FCP	1.2t cDOT
hana			/vol/SS3_log_mnt00001/SS3_log_mnt00001
/dev/sdf	host10	FCP	512.0g cDOT
hana			/vol/SS3_log_mnt00001/SS3_log_mnt00001
/dev/sde	host10	FCP	512.0g cDOT

```
hana /vol/SS3_log_mnt00001/SS3_log_mnt00001
/dev/sdd host10 FCP 512.0g cDOT
hana /vol/SS3_log_mnt00001/SS3_log_mnt00001
/dev/sdc host10 FCP 512.0g cDOT
```

3. Run the `multipath -r` command to get the worldwide identifiers (WWIDs) for the device file names:



In this example, there are four LUNs.

```
stlrx300s8-6:~ # multipath -r
create: 3600a098038304436375d4d442d753878 undef NETAPP,LUN C-Mode
size=512G features='3 pg_init_retries 50 queue_if_no_path' hwhandler='0'
wp=undef
|+- policy='service-time 0' prio=50 status=undef
| |- 10:0:1:0 sdd 8:48 undef ready running
| |- 10:0:3:0 sdf 8:80 undef ready running
| |- 11:0:0:0 sds 65:32 undef ready running
| `-- 11:0:2:0 sdu 65:64 undef ready running
`-+- policy='service-time 0' prio=10 status=undef
| |- 10:0:0:0 sdc 8:32 undef ready running
| |- 10:0:2:0 sde 8:64 undef ready running
| |- 11:0:1:0 sdt 65:48 undef ready running
| `-- 11:0:3:0 sdv 65:80 undef ready running
create: 3600a098038304436375d4d442d753879 undef NETAPP,LUN C-Mode
size=1.2T features='3 pg_init_retries 50 queue_if_no_path' hwhandler='0'
wp=undef
|+- policy='service-time 0' prio=50 status=undef
| |- 10:0:1:1 sdj 8:144 undef ready running
| |- 10:0:3:1 sdp 8:240 undef ready running
| |- 11:0:0:1 sdw 65:96 undef ready running
| `-- 11:0:2:1 sdac 65:192 undef ready running
`-+- policy='service-time 0' prio=10 status=undef
| |- 10:0:0:1 sdg 8:96 undef ready running
| |- 10:0:2:1 sdm 8:192 undef ready running
| |- 11:0:1:1 sdz 65:144 undef ready running
| `-- 11:0:3:1 sdaf 65:240 undef ready running
create: 3600a098038304436392b4d442d6f534f undef NETAPP,LUN C-Mode
size=1.2T features='3 pg_init_retries 50 queue_if_no_path' hwhandler='0'
wp=undef
|+- policy='service-time 0' prio=50 status=undef
| |- 10:0:0:2 sdh 8:112 undef ready running
| |- 10:0:2:2 sdn 8:208 undef ready running
| |- 11:0:1:2 sdaa 65:160 undef ready running
| `-- 11:0:3:2 sdag 66:0 undef ready running
`-+- policy='service-time 0' prio=10 status=undef
```

```

|- 10:0:1:2 sdk 8:160 undef ready running
|- 10:0:3:2 sdq 65:0 undef ready running
|- 11:0:0:2 sdx 65:112 undef ready running
`- 11:0:2:2 sdad 65:208 undef ready running
create: 3600a098038304436392b4d442d6f5350 undef NETAPP,LUN C-Mode
size=512G features='3 pg_init_retries 50 queue_if_no_path' hwhandler='0'
wp=undef
|+-- policy='service-time 0' prio=50 status=undef
| |- 10:0:0:3 sdi 8:128 undef ready running
| |- 10:0:2:3 sdo 8:224 undef ready running
| |- 11:0:1:3 sdab 65:176 undef ready running
| `-- 11:0:3:3 sdah 66:16 undef ready running
`+-- policy='service-time 0' prio=10 status=undef
|- 10:0:1:3 sdl 8:176 undef ready running
|- 10:0:3:3 sdr 65:16 undef ready running
|- 11:0:0:3 sdy 65:128 undef ready running
`- 11:0:2:3 sdae 65:224 undef ready running

```

4. Edit the `/etc/multipath.conf` file and add the WWIDs and alias names.



The example output shows the content of the `/etc/multipath.conf` file, which includes alias names for the four LUNs of a 2+1 multiple-host system. If there is no `multipath.conf` file available, you can create one by running the following command: `multipath -T > /etc/multipath.conf`.

```

stlrx300s8-6:/ # cat /etc/multipath.conf
multipaths {
    multipath {
        wwid      3600a098038304436392b4d442d6f534f
        alias     hana- SS3_data_mnt00001
    }
    multipath {
        wwid      3600a098038304436375d4d442d753879
        alias     hana- SS3_data_mnt00002
    }
    multipath {
        wwid      3600a098038304436375d4d442d753878
        alias     hana- SS3_log_mnt00001
    }
    multipath {
        wwid      3600a098038304436392b4d442d6f5350
        alias     hana- SS3_log_mnt00002
    }
}

```

5. Run the `multipath -r` command to reload the device map.
6. Verify the configuration by running the `multipath -ll` command to list all the LUNs, alias names, and active and standby paths.



The following example output shows the output of a 2+1 multiple-host HANA system with two data and two log LUNs.

```
stlrx300s8-6:~ # multipath -ll
hana- SS3_data_mnt00002 (3600a098038304436375d4d442d753879) dm-1
NETAPP,LUN C-Mode
size=1.2T features='4 queue_if_no_path pg_init_retries 50
retain_attached_hw_handler' hwhandler='1 alua' wp=rw
|+- policy='service-time 0' prio=50 status=enabled
| |- 10:0:1:1 sdj 8:144 active ready running
| |- 10:0:3:1 sdp 8:240 active ready running
| |- 11:0:0:1 sdw 65:96 active ready running
| `-- 11:0:2:1 sdac 65:192 active ready running
`+- policy='service-time 0' prio=10 status=enabled
  |- 10:0:0:1 sdg 8:96 active ready running
  |- 10:0:2:1 sdm 8:192 active ready running
  |- 11:0:1:1 sdz 65:144 active ready running
  `-- 11:0:3:1 sdaf 65:240 active ready running
hana- SS3_data_mnt00001 (3600a098038304436392b4d442d6f534f) dm-2
NETAPP,LUN C-Mode
size=1.2T features='4 queue_if_no_path pg_init_retries 50
retain_attached_hw_handler' hwhandler='1 alua' wp=rw
|+- policy='service-time 0' prio=50 status=enabled
| |- 10:0:0:2 sdh 8:112 active ready running
| |- 10:0:2:2 sdn 8:208 active ready running
| |- 11:0:1:2 sdaa 65:160 active ready running
| `-- 11:0:3:2 sdag 66:0 active ready running
`+- policy='service-time 0' prio=10 status=enabled
  |- 10:0:1:2 sdk 8:160 active ready running
  |- 10:0:3:2 sdq 65:0 active ready running
  |- 11:0:0:2 sdx 65:112 active ready running
  `-- 11:0:2:2 sdad 65:208 active ready running
hana- SS3_log_mnt00002 (3600a098038304436392b4d442d6f5350) dm-3
NETAPP,LUN C-Mode
size=512G features='4 queue_if_no_path pg_init_retries 50
retain_attached_hw_handler' hwhandler='1 alua' wp=rw
|+- policy='service-time 0' prio=50 status=enabled
| |- 10:0:0:3 sdi 8:128 active ready running
| |- 10:0:2:3 sdo 8:224 active ready running
| |- 11:0:1:3 sdab 65:176 active ready running
| `-- 11:0:3:3 sdah 66:16 active ready running
```

```

`-+- policy='service-time 0' prio=10 status=enabled
  |- 10:0:1:3 sdl 8:176 active ready running
  |- 10:0:3:3 sdr 65:16 active ready running
  |- 11:0:0:3 sdy 65:128 active ready running
  `-- 11:0:2:3 sdae 65:224 active ready running
hana- SS3_log_mnt00001 (3600a098038304436375d4d442d753878) dm-0
NETAPP,LUN C-Mode
size=512G features='4 queue_if_no_path pg_init_retries 50
retain_attached_hw_handler' hwhandler='1 alua' wp=rw
|-+- policy='service-time 0' prio=50 status=enabled
| |- 10:0:1:0 sdd 8:48 active ready running
| |- 10:0:3:0 sdf 8:80 active ready running
| |- 11:0:0:0 sds 65:32 active ready running
| `-- 11:0:2:0 sdu 65:64 active ready running
`-+- policy='service-time 0' prio=10 status=enabled
  |- 10:0:0:0 sdc 8:32 active ready running
  |- 10:0:2:0 sde 8:64 active ready running
  |- 11:0:1:0 sdt 65:48 active ready running
  `-- 11:0:3:0 sdv 65:80 active ready running

```

Create LVM volume groups and logical volumes

This step is only needed if LVM will be used. The following example is for a 2+1 host setup using SID FC5.



For an LVM- based setup, the multipath configuration described in the previous section must be completed as well. In this example, eight LUNs must be configured for multipathing.

1. Initialize all LUNs as a physical volume.

```

pvcreate /dev/mapper/hana-FC5_data_mnt00001
pvcreate /dev/mapper/hana-FC5_data2_mnt00001pvcreate /dev/mapper/hana-
FC5_data_mnt00002
pvcreate /dev/mapper/hana-FC5_data2_mnt00002
pvcreate /dev/mapper/hana-FC5_log_mnt00001
pvcreate /dev/mapper/hana-FC5_log2_mnt00001pvcreate /dev/mapper/hana-
FC5_log_mnt00002
pvcreate /dev/mapper/hana-FC5_log2_mnt00002

```

2. Create the volume groups for each data and log partition.


```
vgcreate FC5_data_mnt00001 /dev/mapper/hana-FC5_data_mnt00001
/dev/mapper/hana-FC5_data2_mnt00001
vgcreate FC5_data_mnt00002 /dev/mapper/hana-FC5_data_mnt00002
/dev/mapper/hana-FC5_data2_mnt00002
vgcreate FC5_log_mnt00001 /dev/mapper/hana-FC5_log_mnt00001
/dev/mapper/hana-FC5_log2_mnt00001
vgcreate FC5_log_mnt00002 /dev/mapper/hana-FC5_log_mnt00002
/dev/mapper/hana-FC5_log2_mnt00002
```

3. Create a logical volume for each data and log partition. Use a stripe size that is equal to the number of LUNs used per volume group (in example two) and a stripe size of 256k for data and 64k for log. SAP only supports one logical volume per volume group.

```
lvcreate --extents 100%FREE -i 2 -I 256k --name vol FC5_data_mnt00001
lvcreate --extents 100%FREE -i 2 -I 256k --name vol FC5_data_mnt00002
lvcreate --extents 100%FREE -i 2 -I 64k --name vol FC5_log_mnt00002
lvcreate --extents 100%FREE -i 2 -I 64k --name vol FC5_log_mnt00001
```

4. Scan the physical volumes, volume groups, and vol groups at all other hosts.

```
modprobe dm_modpvscanvgscanlvscan
```



If the above commands do not find the volumes, a restart is required.

5. To mount the logical volumes, the logical volumes must be activated. To activate the volumes, run the following command:

```
vgchange -a y
```

Create file systems

To create the XFS file system on each LUN belonging to the HANA system, take one of the following actions:

- For a single-host system, create the XFS file system on the data, log, and **/hana/shared** LUNs.

```
stlrx300s8-6:/ # mkfs.xfs /dev/mapper/hana- SS3_data_mnt00001
stlrx300s8-6:/ # mkfs.xfs /dev/mapper/hana- SS3_log_mnt00001
stlrx300s8-6:/ # mkfs.xfs /dev/mapper/hana- SS3_shared
```

- For a multiple-host system, create the XFS file system on all data and log LUNs.

```
stlrx300s8-6:~ # mkfs.xfs /dev/mapper/hana- SS3_log_mnt00001
stlrx300s8-6:~ # mkfs.xfs /dev/mapper/hana- SS3_log_mnt00002
stlrx300s8-6:~ # mkfs.xfs /dev/mapper/hana- SS3_data_mnt00001
stlrx300s8-6:~ # mkfs.xfs /dev/mapper/hana- SS3_data_mnt00002
```

- If LVM is used, create the XFS file system on all data and log logical volumes.

```
mkfs.xfs FC5_data_mnt00001-vol
mkfs.xfs FC5_data_mnt00002-vol
mkfs.xfs FC5_log_mnt00001-vol
mkfs.xfs FC5_log_mnt00002-vol
```



The multiple host example commands show a 2+1 multiple-host HANA system.

Create mount points

To create the required mount point directories, take one of the following actions:

- For a single-host system, set permissions and create mount points on the database host.

```
stlrx300s8-6:/ # mkdir -p /hana/data/SS3/mnt00001
stlrx300s8-6:/ # mkdir -p /hana/log/SS3/mnt00001
stlrx300s8-6:/ # mkdir -p /hana/shared
stlrx300s8-6:/ # chmod -R 777 /hana/log/SS3
stlrx300s8-6:/ # chmod -R 777 /hana/data/SS3
stlrx300s8-6:/ # chmod 777 /hana/shared
```

- For a multiple-host system, set permissions and create mount points on all worker and standby hosts.



The example commands show a 2+1 multiple-host HANA system.

```
stlrx300s8-6:/ # mkdir -p /hana/data/SS3/mnt00001
stlrx300s8-6:/ # mkdir -p /hana/log/SS3/mnt00001
stlrx300s8-6:/ # mkdir -p /hana/data/SS3/mnt00002
stlrx300s8-6:/ # mkdir -p /hana/log/SS3/mnt00002
stlrx300s8-6:/ # mkdir -p /hana/shared
stlrx300s8-6:/ # chmod -R 777 /hana/log/SS3
stlrx300s8-6:/ # chmod -R 777 /hana/data/SS3
stlrx300s8-6:/ # chmod 777 /hana/shared
```



The same steps must be executed for a system configuration with Linux LVM.

Mount file systems

To mount file systems during system boot using the `/etc/fstab` configuration file, complete the following steps:

1. Take one of the following actions:

- For a single-host system, add the required file systems to the `/etc/fstab` configuration file.



The XFS file systems for the data and log LUN must be mounted with the `relatime` and `inode64` mount options.

```
stlrx300s8-6:/ # cat /etc/fstab
/dev/mapper/FAS8200-hana- SS3_shared /hana/shared xfs defaults 0 0
/dev/mapper/FAS8200-hana- SS3_log_mnt00001 /hana/log/SS3/mnt00001 xfs
relatime,inode64,nobarrier 0 0
/dev/mapper/FAS8200-hana- SS3_data_mnt00001 /hana/data/SS3/mnt00001
xfs relatime,inode64 0 0
```

If LVM is used, use the logical volume names for data and log.

```
# cat /etc/fstab
/dev/mapper/hana-FC5_shared /hana/shared xfs defaults 0 0
/dev/mapper/FC5_log_mnt00001-vol /hana/log/FC5/mnt00001 xfs
relatime,inode64 0 0
/dev/mapper/FC5_data_mnt00001-vol /hana/data/FC5/mnt00001 xfs
relatime,inode64 0 0
```

- For a multiple-host system, add the `/hana/shared` file system to the `/etc/fstab` configuration file of each host.



All the data and log file systems are mounted through the SAP HANA storage connector.

```
stlrx300s8-6:/ # cat /etc/fstab
<storage-ip>:/hana_shared /hana/shared nfs
rw,vers=3,hard,timeo=600,intr,noatime,nolock 0 0
```

2. To mount the file systems, run the `mount -a` command at each host.

Next: [I/O stack configuration for SAP HANA.](#)

Copyright Information

Copyright © 2021 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system-without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.